

Genome Assembly

officialprofile

2021-09-07

Contents

1	Preface	5
1.1	Prerequisites	6
2	Introduction	7
3	Basic string manipulations	9
3.1	Generate random DNA sequence	9
3.2	Find cDNA	10
3.3	Find reverse complementary	11
3.4	Calculate frequency of the nucleobases	11
3.5	Longest common prefix	12
4	Graph theory	15
5	De Bruijn graph	17
6	Burrows-Wheeler transform	19
6.1	Introduction	19
6.2	Burrows-Wheeler matrix	19
6.3	Inverse transform	21


Chapter 1

Preface

This mini textbook describes selected algorithms that play a main role in the de novo genome assembly. The premise of this book is to construct these algorithms from the very bottom and explain step by step main ideas that stand behind them.

We will try to avoid as much as possible ready-to-use implementations, which are of course available and of great quality, but their use wouldn't serve the educational purpose. Naturally, many applications are included as well.

By default the code is written in R, but at certain points python is also mentioned.



A HiFi De Bruijn graph for a pile of reads from *Drosophila* genome sequencing. Each dot represents a k-mer ($k=23$), the edges denote neighboring k-mers. The larger red dots mark the head of heterozygous bubbles. Source: pacb.com.

1.1 Prerequisites

It is assumed that the reader

1. has a basic understanding of genetics;
2. has some experience with programming in R (is familiar with loops, data structures, etc.)
3. had some contact with higher mathematics, e.g. statistics, graph theory.

Throughout the book the following libraries are being used and it is assumed that the reader has them loaded.

```
library(stringr)
library(dplyr)
library(igraph)
```

Chapter 2

Introduction

Genome assembly has been regarded as one of the most important and most challenging problems in bioinformatics, perhaps at least since the late 80's, when the Human Genome Project was announced.

Genomes tend to be almost inconceivably long. Human DNA for example is approximately thousand times longer than the Bible (letter-wise), and some species have their genomes even orders of magnitude longer. Yet the algorithmic part of the problem arises not mainly because of the length. This was definitely the case for the biotechnologists, who few decades ago were constrained within low-throughput sequencing techniques. For bioinformaticians though the core of the problem is located a bit deeper.

A good analogy that explains the nature of our challenge is a shredded newspaper (some authors call it the newspaper problem, other use book as example). Namely, *de novo* genome assembly resembles reconstruction of the original document from a set of unarranged newspaper pieces, like on the figure below.

As one can imagine the problem is difficult, but not because the newspaper has twenty or forty pages. If we had just one the assembly still wouldn't be easy. The length of the genome, although far from being irrelevant, is for us more of a secondary issue. On the other hand, repetitive patterns, which for biotechnologists are not a problem at all, will complicate our journey substantially. Reader will perhaps come to the very similar conclusions as the succeeding topics will uncover.

One should also underline the fact that transcriptome assembly is not the same as the genome assembly. Read coverage of the latter is relatively uniform, whereas transcriptome can be differentially expressed and therefore the frequency of its reads can vary a lot. In genome assembly non-uniform abundance of the reads simply indicates existence of repetitions. In the case of transcription product it is much less straightforward. Methods that overcome this issue exist but they are beyond of the scope of this textbook.

Chapter 3

Basic string manipulations

Before we dive into the main topics let's warm up with few basic exercises. On the one hand some of these problems might be regarded as a form of general familiarization with string manipulations. Being able to get through them is in a sense a must, at least from a perspective of algorithms that will be discussed later. On the other hand some of the problems, like generating k-mers or suffixes, are absolutely crucial in genome assembly and one has to know them to a tee. So, without further ado, let's do the warm up.

3.1 Generate random DNA sequence

Let's start with writing a function that, for a given integer $n > 0$, returns randomly generated DNA sequence.

```
random_sequence <- function(n){  
  return(sample(c('A','C','G','T'), n, replace = TRUE))  
}
```

```
random_sequence(10)  
#> [1] "G" "C" "T" "A" "T" "G" "G" "A" "C" "T"
```

If we want our returned sequence to be in a form of a single string we can use `paste()` function with `collapse` parameter equal to `' '`.

```
paste(random_sequence(10), collapse = '')  
#> [1] "TTTAGAGAT"
```

Unfortunately, there is no a convenient way of retrieving a substring. If a DNA variable is equal to 'ACGTG' then we won't get a substring 'CG' simply by

writing `DNA[2:3]` (one could do this in python). One could do this by employing the `substr()` function, i.e. `substr(DNA, 2, 3)`. On the other hand, if `DNA` is a vector `c('A', 'C', 'G', 'T', 'G')` then `DNA[2:3]` will return `c('C', 'G')` that in turn can be merged into `'CG'` by `paste()`. Either way, one must use a certain function in order to glue the letters or to split the sequence.

3.2 Find cDNA

To find the complementary sequence one must know that C (cytosine) is complementary to G (guanine), and A (adenine) is complementary to T (thymine). Also, before creating the cDNA we will at first construct a list (in Python it would be a dictionary) of complementary nucleobases.

```
c_bases <- list('A' = 'T', 'C' = 'G', 'G' = 'C', 'T' = 'A')

c_bases['A'][[1]]
#> [1] "T"
```

This data structure allows us to directly access complementary bases. Let's build a function that utilizes this.

```
complementary_sequence <- function(sequence){
  complementary = c()
  for (base in sequence){
    complementary = c(complementary, c_bases[base] [[1]])
  }
  return(complementary)
}
```

Please note that we append complementary nucleobase in a seemingly non-optimal way, namely by writing `sequence = c(sequence, new_nucleobase)`. One could argue that this should be done with `append()`, i. e. `sequence = append(sequence, new_nucleobase)`, but in fact the `append` in R is regarded as a relatively slow function and our solution is usually more recommended (in python using `append` is fine).

```
seq <- random_sequence(20)
cat(' DNA =', seq, '\n')
#> DNA = G C A G C T T C A T T A A A A G A C A
cseq <- complementary_sequence(seq)
cat(' cDNA =', cseq)
#> cDNA = C G T C G A A G T A A T T T T C T G T
```

There is also a neater way to compute, and in a way avoid, these kind of loops. One can achieve this though functions like `apply()`, `sapply()` or `lapply()`.

```
complementary_sequence <- function(sequence){
  sapply(seq, function(base) c_bases[base][[1]])
}
```

The clear downside of using functions like the one above is loss of legibility. Applying them can also be not as straightforward and intuitive as writing a conventional loop. Advanced programmers though find these functions very handy, fast and as readable as standard for or while loop. For that reason we will try to balance these two approaches out. We don't want our code to be too hermetic, but also we should not limit ourselves by deliberately avoiding legitimate solutions.

3.3 Find reverse complementary

Obtaining reverse complementary is very similar. The only difference is that instead of appending complementary nucleobases we will prepend them.

```
reverse_complementary <- function(sequence){
  reverse = c()
  for (base in sequence){
    reverse = c(c_bases[base][[1]], reverse)
  }
  return(reverse)
}
```

```
seq <- random_sequence(20)
cat(' DNA =', seq, '\n')
#> DNA = G G G A G T C T G G T G C A G A C G T
rcseq <- reverse_complementary(seq)
cat(' cDNA =', rcseq)
#> cDNA = A C G T C T G C A C C C A G A C T C C C
```

3.4 Calculate frequency of the nucleobases

```
frequency <- function(sequence, percents = FALSE){
  counts <- table(sequence)
  divide_counts <- ifelse(percents, sum(counts), 1)
  return (counts/divide_counts)
}
```

```
frequency(seq, percents = TRUE)
#> sequence
#>   A   C   G   T
#> 0.15 0.15 0.50 0.20
```

3.5 Longest common prefix

```
longest_common_prefix <- function(sequence1, sequence2){
  min_length = min(length(sequence1), length(sequence2))
  index = 0
  if (sequence1[1] != sequence2[1]){
    return ('No common prefix')
  }
  for (i in 2:min_length){
    if (sequence1[i] != sequence2[i]){
      return (sequence1[1:i-1])
    }
  }
  return (sequence1[1:min_length])
}
```

```
longest_common_prefix(c('A','C','T','G'), c('A', 'C', 'T', 'T', 'T'))
#> [1] "A" "C" "T"
```

```
longest_common_prefix(c('C','C','T','G'), c('A', 'C', 'T', 'T', 'T'))
#> [1] "No common prefix"
```

Inputting strings that are in fact consisted of vectors of characters may not be particularly convenient. We will therefore add a functionality that detects type of inputs and converts them if it is needed.

Unfortunately, R does not differentiate between `c('a','c','t','g')` and `'ACTG'`, e.g.

```
typeof(c('A', 'C', 'G', 'T'))
#> [1] "character"
typeof('ACGT')
#> [1] "character"

length(c('A', 'C', 'G', 'T'))
#> [1] 4
length('ACGT')
#> [1] 1
```

On this account we will simply add new boolean argument `glued` to our function. This solution is far from perfect, but for now is good enough.

```
longest_common_prefix <- function(sequence1, sequence2, glued = FALSE){  
  if (glued){  
    sequence1 <- strsplit(sequence1, '')[[1]]  
    sequence2 <- strsplit(sequence2, '')[[1]]  
  }  
  min_length = min(length(sequence1), length(sequence2))  
  index = 0  
  if (sequence1[1] != sequence2[1]){  
    return ('No common prefix')  
  }  
  for (i in 2:min_length){  
    if (sequence1[i] != sequence2[i]){  
      return (sequence1[1:i-1])  
    }  
  }  
  return (sequence1[1:min_length])  
}
```

```
longest_common_prefix('ACTG', 'ACCC', glued = TRUE)  
#> [1] "A" "C"
```


Chapter 4

Graph theory

Chapter 5

De Bruijn graph

Chapter 6

Burrows-Wheeler transform

The Burrows-Wheeler transform is one of the most effective lossless text compression method available. It provides a reversible transformation for text that makes it easier to compress. Of course, one may wonder what text compression has to do with genome assembly. As a matter of fact these two issues are closely related. But we should be more precise here - text compression is closely related to pattern matching which in turn is crucial for the genome assembly. In a broad sense compression algorithms look for patterns and try to remove repetitions. We want to take advantage of this feature, especially because repetitive patterns tend to be very abundant in genomic sequences.

It is worth mentioning that the Burrows-Wheeler transform is also closely related to suffix trees and suffix arrays, which are commonly used within pattern matching. This relationship will be studied later but perhaps reader should already keep the trivia in mind. (Donald Adjero, 2008)

6.1 Introduction

The Burrows-Wheeler transform method is often referred to as “block sorting”, because it takes a block of text and permutes it. By permuting a block of text we mean rearranging the order of its symbols. Once again, we should be more precise here because Burrows-Wheeler transform performs a specific type of permutation, namely *circular shift permutation*: all of the characters are moved one position to the left, and first character moves to the last position.

6.2 Burrows-Wheeler matrix

Consider the following sequence:

```
1 sequence <- 'GATTACA'
```

In order to create the Burrows-Wheeler matrix, from which the transform itself can be obtained, for the given string we at first add the dollar sign \$ at the end of the sequence.

```
1 sequence <- str_c(sequence, '$')
```

Afterwards we perform a series of circular shift permutations.

```
1 sequences <- c(sequence)
2 n          <- nchar(sequence)
3
4 for (i in 1:(n-1)){
5   sequence <- str_c(str_sub(sequence, 2, n),
6                     str_sub(sequence, 1, 1))
7
8   sequences <- c(sequences, sequence)
9 }
10
11 cat(sequences, sep = '\n')
12 #> GATTACA$
13 #> ATTACA$G
14 #> TTACA$GA
15 #> TACA$GAT
16 #> ACA$GATT
17 #> CA$GATTA
18 #> A$GATTAC
19 #> $GATTACA
```

Then we sort these sequences with the assumption that the dollar sign precedes lexicographically every other symbol.

```
1 sequences <- sort(sequences)
2 cat(sequences, sep = '\n')
3 #> $GATTACA
4 #> A$GATTAC
5 #> ACA$GATT
6 #> ATTACA$G
7 #> CA$GATTA
8 #> GATTACA$
9 #> TACA$GAT
10 #> TTACA$GA
```

For our convenience let's split these permutations into vectors of single characters.

```

1 bw.matrix      <- data.frame(matrix(, n, n))
2 colnames(bw.matrix) <- 1:n
3
4 for (i in 1:n){
5   bw.matrix[i, ] <- strsplit(sequences[i], split = '')[[1]]
6 }
7
8 knitr::kable(bw.matrix)

```

1	2	3	4	5	6	7	8
\$	G	A	T	T	A	C	A
A	\$	G	A	T	T	A	C
A	C	A	\$	G	A	T	T
A	T	T	A	C	A	\$	G
C	A	\$	G	A	T	T	A
G	A	T	T	A	C	A	\$
T	A	C	A	\$	G	A	T
T	T	A	C	A	\$	G	A

Thus we have created the **Burrows-Wheeler matrix**. Sequence in the last column is called the **Burrows-Wheeler transform**.

```

1 transform <- paste(bw.matrix[,n], collapse = '')
2
3 cat('The Burrows-Wheeler transform of',
4     sequence, 'is', transform)
5 #> The Burrows-Wheeler transform of $GATTACA is ACTGA$TA

```

6.3 Inverse transform

As we said at the very beginning the transform is reversible. Having only the transformed sequence we are going to reconstruct the Burrows-Wheeler matrix and initial sequence itself.

Firstly let's sort the characters of the transformed sequence.

```

1 first.sequence <- strsplit(transform, split = '')[[1]] %>% sort
2 paste(first.sequence, collapse = '')
3 #> [1] "$AAACGTT"

```

Note that this string is equivalent to the first column of the Burrows-Wheeler transform.

```

1 bw.inverse      <- data.frame(matrix(, n, 2))
2 colnames(bw.inverse) <- c(n, 1)
3
4 bw.inverse[, 1] <- strsplit(transform, split = '')[[1]]
5 bw.inverse[, 2] <- first.sequence
6
7 knitr::kable(bw.inverse)

```

8	1
A	\$
C	A
T	A
G	A
A	C
\$	G
T	T
A	T

Also keep in mind that the characters from last and the first column are adjacent. In other words, at this point we have a set of 2-mers.

```

1 kmers <- apply(bw.inverse, 1,
2               function(x) paste(x, collapse = ''))
3 kmers
4 #> [1] "A$" "CA" "TA" "GA" "AC" "$G" "TT" "AT"

```

The reconstruction process strictly relies on the fact that Burrows-Wheeler matrix is sorted lexicographically. This property will allow us to retrieve the remaining columns.

```

1 kmers <- sort(kmers)
2 kmers
3 #> [1] "$G" "A$" "AC" "AT" "CA" "GA" "TA" "TT"

```

The 2-mers (k-mers in general) that we sorted lexicographically represent first two columns of the Burrows-Wheeler matrix. We can extract last character of each 2-mer in the following way:

```

1 sapply(kmers, function(x) str_sub(x, 2, 2),
2       simplify = TRUE, USE.NAMES = FALSE)
3 #> [1] "G" "$" "C" "T" "A" "A" "A" "T"

```

By inserting this set of characters we obtained the second column, and by iterating the process of building substrings, sorting them, and retrieving last characters we can fill the whole Burrows-Wheeler matrix.

```

1 for (i in 2:(n-1)){
2   kmers      <- apply(bw.inverse, 1,
3                       function(x) paste(x, collapse = ''))
4   kmers      <- sort(kmers)
5   bw.inverse[, i+1] <- sapply(kmers, function(x) str_sub(x, i, i),
6                               simplify = TRUE, USE.NAMES = FALSE)
7   colnames(bw.inverse)[i+1] = i
8 }
9 knitr::kable(bw.inverse)

```

8	1	2	3	4	5	6	7
A	\$	G	A	T	T	A	C
C	A	\$	G	A	T	T	A
T	A	C	A	\$	G	A	T
G	A	T	T	A	C	A	\$
A	C	A	\$	G	A	T	T
\$	G	A	T	T	A	C	A
T	T	A	C	A	\$	G	A
A	T	T	A	C	A	\$	G

Finally we move first column to the very end

```

1 bw.inverse[,n+1] <- bw.inverse[, 1]
2 bw.inverse      <- bw.inverse[,2:(n+1)]
3 colnames(bw.inverse)[n] = n
4
5 knitr::kable(bw.inverse)

```

1	2	3	4	5	6	7	8
\$	G	A	T	T	A	C	A
A	\$	G	A	T	T	A	C
A	C	A	\$	G	A	T	T
A	T	T	A	C	A	\$	G
C	A	\$	G	A	T	T	A
G	A	T	T	A	C	A	\$
T	A	C	A	\$	G	A	T
T	T	A	C	A	\$	G	A

One can also verify that `bw.matrix` and `bw.inverse` are in fact the same.

```

1 knitr::kable(bw.inverse == bw.matrix)

```

1	2	3	4	5	6	7	8
TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE

Additionally we can encapsulate the Burrows-Wheeler transform in a form of a single function.

```

1 BWT <- function(sequence){
2   sequence <- str_c(sequence, '$')
3   sequences <- c(sequence)
4   n <- nchar(sequence)
5
6   for (i in 1:(n-1)){
7     sequence <- str_c(str_sub(sequence, 2, n),
8                       str_sub(sequence, 1, 1))
9     sequences <- c(sequences, sequence)
10  }
11  sequences <- sort(sequences)
12
13  bw.matrix <- data.frame(matrix(, n, n))
14  colnames(bw.matrix) <- 1:n
15
16  for (i in 1:n){
17    bw.matrix[i, ] <- strsplit(sequences[i], split = '')[[1]]
18  }
19  return(paste(bw.matrix[,n], collapse = ''))
20 }
```

```

1 BWT('GATTACA')
2 #> [1] "ACTGA$TA"
```

One can verify that this output is equal to result we obtained earlier.

Out of pure curiosity lets check the Burrows-Wheeler transform for a longer sequence.

```

1 BWT('ATGCTCGTGCCATCATATAGCGCGCGCGATCTCTACGCGCG')
2 #> [1] "GTTTCCG$TCGGGGGAGGGTTGTCTCCCCCATCAAACCAGA"
```

Please note that the input string has no identical characters at adjacent positions, whereas in the transformed sequence such situation appears quite often.

These substrings of identical characters will allow us represent the sequence in a more condensed manner and expediate pattern matching.

Bibliography

Donald Adjero, Tim Bell, A. M. (2008). *The Burrows-Wheeler Transform: Data Compression, Suffix Arrays, and Pattern Matching*. Springer.